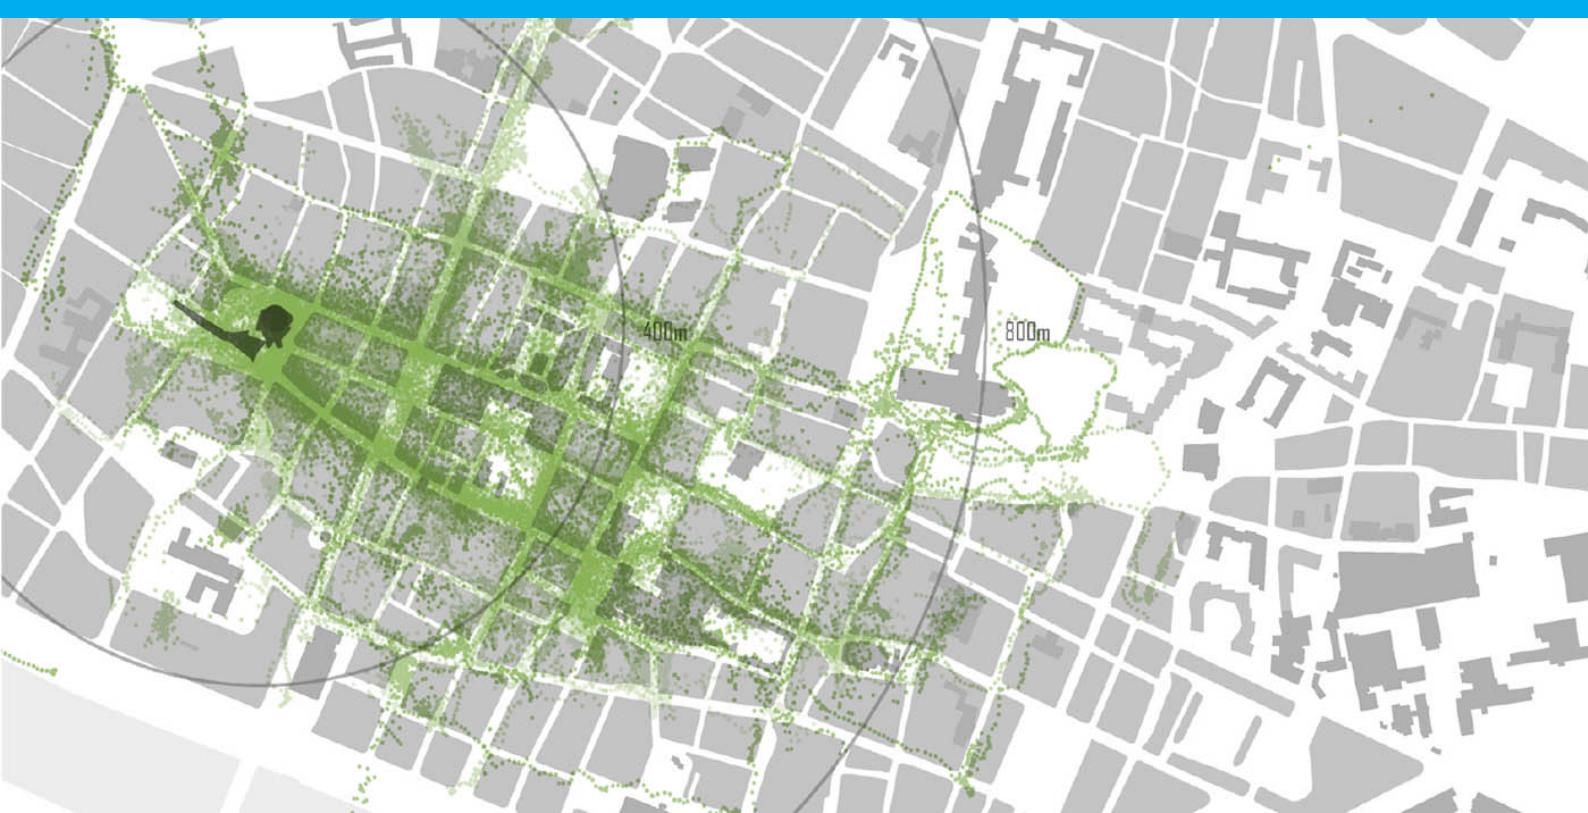


Identifying movement patterns from large scale WiFi-based location data

A case study of the TU Delft Campus

Balazs Dukai
Simon Griffioen
Matthijs Bon
Martijn Vermeer
Xander den Duijn
Yuxuan Kang



Identifying movement patterns from large scale WiFi-based location data

A case study of the TU Delft Campus

by

Balazs Dukai
Simon Griffioen
Matthijs Bon
Martijn Vermeer
Xander den Duijn
Yuxuan Kang

**Synthesis project of Geomatics
at the Delft University of Technology,**

Project duration: April 16, 2016 – June 17, 2016

to be presented on Friday May 20, 2016.



Preface

*Balazs Dukai
Simon Griffioen
Matthijs Bon
Martijn Vermeer
Xander den Duijn
Yuxuan Kang
Delft, May 2016*

Contents

1 Preface	1
2 Glossary	3
3 Executive Summary	5
4 Introduction	7
4.1 Intro	7
4.2 Purpose statement	8
4.3 Methods	8
4.4 Top level requirements	9
4.5 Reading guide.	10
5 Context	11
5.1 Use case: TU Delft (working title)	11
5.2 Previous research: Rhythm of the campus	11
5.3 Privacy	12
5.4 Data accuracy.	12
5.5 Representativeness	12
5.6 Data description and System of APs.	12
5.6.1 Data description	12
5.6.2 System of APs	15
6 Movement patterns	17
6.1 Introduction	17
6.2 Movement identification	17
6.2.1 Spatio-temporal movement patterns.	17
6.2.2 Co-location in space	18
7 Methodology	19
7.1 Preprocessing.	20
7.1.1 Initial filtering	20
7.1.2 Grouping of states	21
7.1.3 Adding world state.	21
7.1.4 Passing by	22
7.1.5 Implementation	23
7.1.6 Apname vs maploc.	24
7.1.7 Static and mobile devices	25
7.2 States to movements	26
8 Spatio-temporal movement patterns	27
8.1 Introduction	27
8.2 Methods	27
8.3 Results	27
8.3.1 All movement	27
8.3.2 Mobile vs static	27
8.3.3 Week vs weekend	27
8.3.4 From and to	27

9 Trajectory patterns	29
9.1 Introduction	29
9.2 Problem description	29
9.2.1 Location extraction	29
9.2.2 Individual trajectory	30
9.2.3 Trajectory Pattern	30
9.3 Implementation	30
9.4 Results	31
10 Indoor spatio-temporal movement patterns	33
10.1 Introduction	33
10.2 Theory / methods.	33
10.3 Implementation	34
10.3.1 Delineation of building parts.	34
10.3.2 Movement between building parts.	35
10.3.3 Indoor route network	35
10.3.4 Mapping traffic to the route network.	35
10.3.5 Visualization of the movement.	36
10.4 Results	36
11 Conclusions	39
12 Recommendations	41
12.1 Entrances	41
12.1.1 Introduction	41
12.1.2 Methodology.	41
12.1.3 Hypothesis.	41
12.1.4 First approach: including passing by events	41
12.1.5 Second approach: excluding passing by events	42
12.1.6 Recommendation	44
12.2 Association rules	45
12.3 Distinguishing user groups	45
12.4 Occupancy	45
12.5 AP system	45
12.6 Data reasoning	45
12.7 Visual exploration.	46
13 Acknowledgements	47
14 Appendix B	49

1

Preface

During the fourth quarter of the first year of the MSc Programme Geomatics for the Built Environment at the TU Delft, the Geomatics Synthesis Project (GSP) takes place. This report is part of this framework and in this project, students will apply all their knowledge they have acquired during the courses while working in groups of five or six students. The students will gain experience throughout the entire process of project management, data processing, data analysis, application and presentation.

This year, the GSP focusses on Wi-Fi tracking data from the eduroam network of the TU Delft. The student will be divided into three groups, each researching one of three different topics:

- Identifying occupancy
- Identifying movement patterns
- Identifying activities

This project is dedicated to the second topic, identifying movement patterns. The project requires 3 main documents: **1)** the baseline review; **2)** the mid term review, and **3)** the final review. This document embodies the final review and was created to provide the students, the supervisor(s) and other involved parties with an overview of the project. The document includes the problem description, development process, results, conclusions and recommendations for future work.

Delft, University of Technology
June, 2016

2

Glossary

Used terms and abbreviations:

Faculty names

AE / LR	Aerospace Engineering
BK / BK City	Faculty of Architecture
CiTG	Civil Engineering
EGM	Thermal Power Plant
EWI / EEMCS	Faculty of Electrical Engineering, Mathematics and Computer Science
FMVG / FMRE	Facility Management & Real Estate
HSL	Hypersonic Wind Tunnel
ID / IO	Industrial Desgin
ISD	International School Delft
LMS	Logistics and Environmental Services
LSL	Low Turbulence Tunnel
O&S	Onderwijs & Studentenzaken
RID	Reactor Institute Delft
SC	Sport Center
TNW	Applied Sciences
TPM	Technology, Policy and Management

Abbreviations

AP / APs	Access Point(s)
GNSS	Global Navigation Satellite System
GSP	Geomatics Synthesis Project
RFID	Radio frequency identification
RSSI	Received Signal Strength Indicator
SNR	Signal to Noise Ratio
SQL	Structured Query Language
WLAN	Wireless Local Area Network

Commonly used terms

Eduroam	Wireless network available at the TU Delft, used interenationally.
Movement	A movement is always from the location of one state to the location of another state, where two states can not be the same.
Pattern	Recurring event that helps in the identification of phenomena.
Region / Buildingpart	Buildingparts and regions refer to large indoor areas that can be grouped together, i.e. 'Staff area', 'Atelier'.
Sequence	Ordered collection of states.
Spatial level	A spatial level defines the level on which states are aggregated
States / Stay places	A state is defined as a time interval during which a particular device is located in a certain area.
Trajectory	A trajectory is defined for each person by an ordered list of buildings that were visited.
World	Location which, depending on the spatial level, can be either outside a building or outside the campus

3

Executive Summary

4

Introduction

4.1. Intro

Wireless Local Area Networks (WLAN) are widely used for indoor positioning of mobile devices within this network. The use of the Wi-Fi network to estimate the location of people is an attractive approach, since Wi-Fi access points (AP) are often available in indoor environments. Furthermore, smart phones are becoming essential in daily life, making it convincing to track mobile devices. This provides a platform to track people by using Wi-Fi monitoring technology. Knowledge of people's locations and related routine activities are important for numerous activities, such as urban planning, emergency rescue and management of buildings.

To understand the human motion behaviour many studies are conducted based on data collection of GPS receivers. The Global Navigation Satellite System (GNSS) is commonly used to track people in large scale environments. However due to poor quality of received signals from satellites in urban or indoor environments, GNSS receivers are not suitable in these environments. Moreover, GNSS receivers are convenient for self-tracking, but for large scale movement analysis, this data should be made available first before others can use it. This led to the development of alternative technologies to track people's locations, including Bluetooth, Dead Reckoning, Radio frequency identification (RFID), ultra-wideband (UWB) and WLAN (Mautz 2012). WLAN has the advantage of widespread deployment, low cost and with the use of a smartphone as a receiver, the possibility to track a large amount of people.

In general, there are four different location tracking techniques by using the Wi-Fi network: Propagation modelling, multilateration, Fingerprinting and Cell of Origin (CoO). Many of these methods rely on Received Signal Strength Indicators (RSSI) and/or previous set of calibration measurements. In comparison, CoO is the most straightforward technique and uses the location of the AP, to locate the mobile device. For, the location of the AP a mobile device is connected to, will give an estimation of the mobile devices' location, and thus the person. For this project, APs related to buildings and building-parts are used to track people's movement.

At the Technical University of Delft (TU Delft) a large scale Wi-Fi network is deployed across all facilities covering the indoor space of the campus. The network is known as an international roaming service for users in educational environments and called the eduroam network. It allows students and staff members from one university to use the infrastructure throughout the campus for free. This allows for large scale collection of Wi-Fi logs including individual scans of mobile devices. A continuous collection of re-locations of devices to access points for a long duration will return detailed records of people's movement. This ubiquitous and individual history location data derived from smartphones will present valuable knowledge on movement on the campus. For this reason, the project is carried out in request of the University's department of Facility Management and Real Estate (FMRE).

In this project, Wi-Fi monitoring technology is used to discover movement patterns on the campus of TU Delft. Based on the relationship between activities and places, location history can be used to discover significant places, movement patterns and hotspots. FMRE can use this information to answer questions such "what is the relation between buildings", "where do people come from" and "how regular a trajectory occurs".

This project will present a method for identification of movement patterns in a large scale indoor environments and between buildings. The method uses concepts of sequential pattern mining. Previous research has been done on sequential pattern mining, such as Zhao et al. 2014 to discover people's life patterns from mobile Wi-Fi scans, Meneses and Moreira 2012 analysed place connectivity using the eduroam network and Radaelli et al. 2013 identifies indoor movement patterns by analysing a sequence of relocations. Individual movement can be identified as a sequence of relocations of a mobile device to different APs. Without any data between two subsequent re-locations, sequential analysis is a convincing way for identifying moving patterns from wifilogs.

4.2. Purpose statement

The project is initiated by the idea that communication technologies can also be used to collect information about connections and connection attempts to Access Points (APs). This geo-referenced information can potentially be used to: **1)** estimate the number of devices at a location at a certain time, representing presence of people at that place at that time or for a certain duration; **2)** track unique ID's over several APs, reconstructing individual patterns of movement, resulting in aggregated flows of people and; **3)** define regular and irregular (temporal, deviating) activities at specific places.

This research will focus on the second matter. Identifying movement patterns has attracted significant interest in recent years. Numerous methods have been explored including Wi-Fi tracking. This report will explain how movement patterns can be identified using large scale Wi-Fi based location data, and tries to contribute with four proposes. **1)** A method for identifying movement patterns by analysing individual sequences of relocations from a large scale Wi-Fi network; This includes filtering the raw data and automatically create individual trajectories over a time interval as a sequence of relocations; **2)** Identify spatio-temporal movement patterns of large crowds of people; **3)** Investigate different visualization methods for showing movement, based on a large scale Wi-Fi network. **4)** A method for analysing indoor movement using a constructed network graph of the underlying building floorplan.

The contributions can be described in one research question for this project.

- How can movement patterns be identified from large scale Wi-Fi-based location data of the eduroam network?

In order to answer the research question, there are three applied subquestions:

- What patterns can be identified moving from and to the TU Delft campus?
- What movement patterns can be identified between buildings on TU Delft campus?
- What movement patterns can be identified between large indoor regions of the Faculty of Architecture?

Besides looking at this project from a spatial pattern perspective, this research also aims to investigate the following topics:

- Privacy – how viable is the data for personal concerns?
- Validity & Accuracy – how reliable is the data, how accurate, how robust for errors?
- Representativeness – which amount of the actual users is covered? Is this ratio constant or location dependant?
- System of APs – how well is the system equipped for measuring and tracking, and what is missing /essential to use the system this way?

4.3. Methods

The Geomatics Synthesis Project (GSP) is a small research project that combines a literature study with practical research. This includes a case study of the TU Delft campus, using real-world data. Practical work includes data storing, processing, analysing, interpretation, visualization and validation. The project is carried out in a team of six students with a connection to a supervisor and stakeholders (FMRE). This involves interactive discussions between stakeholders as an important part of the research.

4.4. Top level requirements

To keep track of the progress of the project, it is necessary to monitor to which degree the project is meeting the top level requirements and if the project is still on schedule with these requirements. In the baseline review the requirements are specified using the MoSCoW rules and killer requirements. In this chapter these requirements will be discussed.

MUST building level

- Main goal: Identify movement patterns and connectivity between building entrances.
- Relate entrances (place) of buildings to the corresponding APs (location).
- Find the stay places of each individual in order of the scan time.
- Find individual trajectories from a sequence of stay places.
- Find the movement patterns, by deriving a sequence of common places shared by all trajectories.
- Visualize the movement patterns between buildings in static maps.

A killer requirement for this level is:

- Identification of APs relating to an entrance of a building

SHOULD buildingpart level

- Main goal: Identify movement patterns between large indoor regions.
- Create a network graph from the underlying building floorplan for the analysis, where each region is a node.
- Find the movement trajectories between regions as a sequence of stays.
- Find the movement patterns between large indoor regions.
- Visualize the movement patterns between regions of buildings.

The killer requirements for this level are:

- Digital indoor floorplan of the buildings with classified/named regions (e.g. study rooms, canteen, etc.)
- Georeferenced building floorplans with APs.

COULD room level

- Main goal: Classification of movement patterns at room level.

The killer requirements for this level are:

- Digital indoor floorplan of the buildings with classified/named rooms (offices, classrooms, project studios, corridors, etc)
- Location of access points
- Fingerprinting map

The following chapters will reflect on these requirements, indicating how successful the project is.

4.5. Reading guide

5

Context

5.1. Use case: TU Delft (working title)

This project's main area of interest is the campus of the TU Delft. There are more than 20,000 students using the campus on more than 150 hectares. This emphasizes even more the magnitude of this project. The network logs the devices connected to the eduroam access points, which implicitly means logging the (approximate) location of the person carrying the device and more information. This tracking data can be used to derive information about the personality of the person carrying the device, such as the distinction between staff and students, based on the tracked locations. Connection to the Wi-Fi eduroam network is free of charge and requires only a NetID, which all students and staff get upon registration at the university.

It is very important to understand, that 'no data is also data'. This means that a device that is not being tracked by any access point for a period of time, is either off-campus or disconnected and still on campus. This provides valuable information when researching the movement patterns. This will be further discussed in the section 7.1.

The eduroam network of the TU Delft campus consists of 1730 access points, distributed over more than 30 buildings. The data is collected for each of the access points over a period of little more than 3 months. The logs are stored in a database on a virtual server, where it is accessible to the three project groups and the Geomatics staff. The data that is collected and the storage in the database is further described in subsection 5.6.1.

The department of Facility Management and Real Estate (FMRE) is the main client for the entire Synthesis Project. They would like to know how the campus is being used, what the hotspots on campus and in buildings are, when people travel the most from one building to another and which buildings are most visited.

5.2. Previous research: Rhythm of the campus

In the fall of 2014, similar research was conducted during another edition of the Geomatics Synthesis Project. The group "Rhythm of the campus" investigated the use of the Library and the Aula of the TU Delft, to gain insight in patterns the use of the facilities of the Library and Aula. This section will give a short summary of their research (Van der Ham et al. 2014).

During the project, the group used passive Wi-Fi monitoring to detect users of the TU Delft Library and the Aula to gain insight in the occupation, in request of FMRE. They used BlueMark sensors at the Library, Aula and 5 other faculties for a period of one week and collected ground truth data for 2 days. Due to its sheer size, the raw data was difficult to process. The data was filtered from static devices and outliers and the data analysis resulted in identification of the occupation of the Library and the Aula. The end results was a dashboard which visualized the sensor network, data analysis and pattern recognition to help the client in the decision making process.

This research was different from the research conducted in this Synthesis Project, mainly due the larger size of the eduroam network and the ability to track everybody using the Wi-Fi network.

5.3. Privacy

This project focuses on identifying common movement patterns, ignoring the individual, therefore we did not test explicitly whether it is possible to identify individuals or not from the data. However, based on our findings about the operation of the *eduroam* Wi-Fi network and about the methods that are used to identify movement patterns, we can make the following assumptions.

Movement patterns are rather unique, therefore it is possible to match them to individuals even if maybe not in every case. However, in order to do so it is necessary to have additional data available. This additional data itself is often considered private data, e.g. the complete weekly schedule of the person. Provided that timetables are openly accessible and the occupation of the individual is known, then his movement pattern may be identified in the dataset.

The availability of a detailed access point map makes it easier to identify individuals by allowing a more detailed movement analysis (e.g. on buildingpart level). It reduces the ambiguity that is still present in building level movement analysis.

5.4. Data accuracy

The spatial accuracy of the Wi-Fi log dataset is defined by the range of the APs. Although we do not have information on the exact range of the different APs, we estimate the range to be a few tens of meters. Therefore, if a user is recorded by a specific AP, in reality he can be anywhere around the AP in its range.

The temporal accuracy of the Wi-Fi log dataset is defined by the five minute campus-wide scan interval of the *eduroam* system. It means that all APs on the TU Delft campus scan at the same moment in approximately five minute intervals. Therefore, it is possible that the user is already at a given AP, but he will be first recorded at the next scan round, or the user already left the AP but that also will be only recorded at the next scan round.

5.5. Representativeness

In the GSP a big amount of wifilog data is used. The data represents all people that make (active) use of the wifi eduroam network. These are the students and employees of the TU Delft. There is just a small amount of people that are within the spatial scope of the project and cannot connect with the wifi eduroam network. The data is acquired by the access points, which all are located in a building on the campus. The people that use a building on the campus, but do not make use of the wifi eduroam network, is very small part. Thus, the main part of actual users is covered by the data used in the GSP. The collection of data is acquired over a continuous time interval of more than 2 months. This time period would be large enough to reflect on all users of the campus to some extend.

5.6. Data description and System of APs

5.6.1. Data description

This section will describe the main datasource within the Synthesis Project; a PostgreSQL database containing the logs from the Wi-Fi scanners on the TU Delft campus. Each row in the wifilog table provides a data value for each column (Table 5.1).

username	mac	asstime	apname	maploc	sesdur	snr	ssi
j85cCQ..	l6iOu+..	14-4-2016 12:30	A-23-0-029	..CITG >4e Verdieping	1:32:02	35	-57
wrBqM..	f2Pw/P..	14-4-2016 7:49	A-23-0-035	..CITG >5e & 6e Verdieping	5:32:16	37	-56
wrBqM..	f2Pw/P..	14-4-2016 13:22	A-23-0-035	..CITG >5e & 6e Verdieping	0:40:20	46	-50
wrBqM..	f2Pw/P..	14-4-2016 14:02	A-23-0-093	..CITG >5e & 6e Verdieping	1:27:13	11	-86
wrBqM..	f2Pw/P..	14-4-2016 15:29	A-23-0-091	..CITG >5e & 6e Verdieping	0:05:08	30	-65
wrBqM..	f2Pw/P..	14-4-2016 15:34	A-23-0-035	..CITG >5e & 6e Verdieping	1:42:32	29	-65
J0IwA+..	HkLY1U..	14-4-2016 11:33	A-23-0-035	..CITG >5e & 6e Verdieping	1:27:40	33	-59
J0IwA+..	HkLY1U..	14-4-2016 13:01	A-23-0-035	..CITG >5e & 6e Verdieping	1:01:01	26	-68
J0IwA+..	HkLY1U..	14-4-2016 14:02	A-23-0-035	..CITG >5e & 6e Verdieping	3:30:19	25	-68
J0IwA+..	HkLY1U..	14-4-2016 17:32	A-23-0-035	..CITG >5e & 6e Verdieping	0:40:05	27	-69

Table 5.1: A segment of the main datasource; the wifilog table

The data value for each attribute (column) in the wifilog table will be described in more detail.

Username

The username column provides the username, as a hashed text. Every user has a unique username, but can appear in the data more than once.

Mac

The mac column provides the media access control address (MAC address), as a hashed text. The MAC address is a unique identifier assigned to a specific piece of hardware, such as the network adapter located in Wi-Fi devices (mobile phones, tablets, laptops etc.). So, it would be possible that a user can have more than one device connected to the Wi-Fi eduroam network.

Asstime

The asstime is the time of which a connected device is recorded by the system.

Apname

The apname is the name assigned to the access point. Every access point has a unique name.

Maploc

The maploc describes the location of the access point. There could be multiple access points with the same maploc. For instance, there are 31 access points located on the ground floor of the Faculty of Architecture.

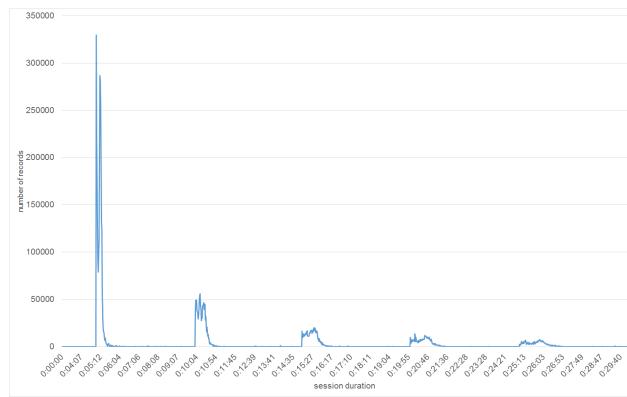


Figure 5.1: The frequency of session durations

Sesdur

The sesdur describes the session duration of which a device is connected to the access point. Because this is not as straightforward as it seems, this will be explained more extensively. Figure 5.1 shows the frequency of session durations (the peak at exactly 5 minutes is filtered out to make the graph more readable). There is a

large peak at exactly 5 minutes, a peak at approximately 5 minutes and decreasing peaks after a time interval of approximately 5 minutes. It looks like it is recording in a certain time interval in which the device is (still) connected.

In order to justify this, the query below is used to see the asstimes (and time to next asstime)

```
select *, asstime_next-asstime as difference
from (
    select count(*), asstime, lead(asstime) over (order by asstime) asstime_next
    from wifilog
    where extract(day from asstime) = 4
    and extract(month from asstime) = 4
    and extract(year from asstime) = 2016
    group by asstime
    order by asstime) as subquery
```

count	asstime	asstime_next	difference
2578	4-4-2016 13:04	4-4-2016 13:09	0:05:10
2435	4-4-2016 13:09	4-4-2016 13:15	0:05:11
2486	4-4-2016 13:15	4-4-2016 13:20	0:05:11
2530	4-4-2016 13:20	4-4-2016 13:25	0:05:11
2471	4-4-2016 13:25	4-4-2016 13:30	0:05:11
2444	4-4-2016 13:30	4-4-2016 13:35	0:05:11
2524	4-4-2016 13:35	4-4-2016 13:40	0:05:11
2588	4-4-2016 13:40	4-4-2016 13:46	0:05:12
2690	4-4-2016 13:46	4-4-2016 13:51	0:05:11
2560	4-4-2016 13:51	4-4-2016 13:56	0:05:11

Table 5.2: The time and time to next scan at a random day

Table 5.2 shows that the time to the next scan is 5 minutes and several seconds in all cases. Most important is to know that all access points are recording the connected device(s) is at the same time.

Table 5.3 will be used to explain the way the time interval of approximately 5 minutes is coming back in the session duration.

The first record shows the device is not connected to any of the access points on the campus in the subsequent moment of recording, resulting in a session duration of exactly 5 minutes. In the last record in Table 5.3 shows the result of a device that is still connected to the same access point at the subsequent moment of recording. In this case the session duration will be 10 minutes and 21 seconds. This is the time interval between the first moment the device is recorded and the first time the device is not recorded by the same access point anymore. The record with id number 6 describes a situation in which the device is connected to an access point at the moment of recording and connected to another access point at the subsequent moment of recording, the session duration is 5 minutes and 18 seconds in this case. This is the time interval between the two moments of recording.

id	username	mac	asstime	apname	maploc	sesdur
1	oHh0Sz..	WWW0Cd..	1-4-2016 10:13	A-12-0-104	..& Proeffabriek >le Verdieping	0:05:00
2	oHh0Sz..	WWW0Cd..	1-4-2016 10:18	A-132-0-064	..32-OCP-IO >1e Verdieping	0:20:27
3	oHh0Sz..	WWW0Cd..	1-4-2016 11:36	A-132-0-105	Root Area	0:15:22
4	oHh0Sz..	WWW0Cd..	1-4-2016 11:51	A-132-0-066	..32-OCP-IO >1e Verdieping	0:20:35
5	oHh0Sz..	WWW0Cd..	1-4-2016 14:01	A-132-0-069	..32-OCP-IO >1e Verdieping	0:05:43
6	oHh0Sz..	WWW0Cd..	1-4-2016 14:06	A-132-0-133	..32-OCP-IO >4e Verdieping	0:05:18
7	oHh0Sz..	WWW0Cd..	1-4-2016 14:12	A-132-0-066	..32-OCP-IO >1e Verdieping	0:05:10
8	oHh0Sz..	WWW0Cd..	1-4-2016 14:17	A-132-0-104	..32-OCP-IO >2e Verdieping	0:05:10
9	oHh0Sz..	WWW0Cd..	1-4-2016 14:22	A-132-0-067	..32-OCP-IO >1e Verdieping	0:05:10
10	oHh0Sz..	WWW0Cd..	1-4-2016 14:27	A-132-0-066	..32-OCP-IO >1e Verdieping	0:10:21

Table 5.3: Varying session durations

SNR

The signal to noise ratio(SNR) describes a measurement that compares the signal strength to the level of background noise (in dB).

RSSI

The received signal strength indicator (RSSI) describes the received signal strength (in dB).

5.6.2. System of APs

This section will describe the current layout of access points (APs) on the TU Delft campus. The location of APs in a building is not known, but for the Faculty of Architecture a paper map was available. Therefore the system of APs in the Faculty of Architecture will be described in more detail.

In total there are 1730 access points, distributed over more than 30 buildings on the campus. The access points are mostly placed on walls or ceilings. The data describes that every access point is linked to a certain location. Due to the (wide) signal range of the access point, the device can be located at a different floor level than the access point it is connected to. Moreover, there could be access points located at the first floor while serving people at ground floor as well. This is the case in rooms with high ceilings, such as the orange hall in the Faculty of Architecture.

As said, the Faculty of Architecture is the only building of which the location of the access points are known. The floor plans are enriched with the location of the access point (see). Next to that, a table is provided with additional information regarding the access points, although this table does not contain all present access points. This table includes the MAC address of the access point. This could be used to look up to what access point the device is connected.

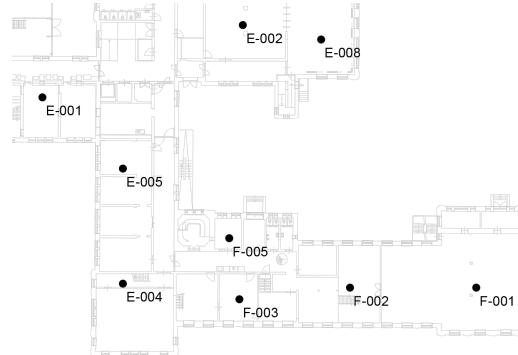


Figure 5.2: Ground floor plan with the location of the access points

6

Movement patterns

6.1. Introduction

The objective of this project is to identify movement patterns. To have a better understanding of this concept, it is important to describe relevant types of movement patterns in a systematic and comprehensive way. A classification of different patterns will provide guidelines for development of different mining algorithms and identify patterns. This chapter will first approach the definition of movement patterns. Subsequently, the theory is demonstrated with the research case of TU Delft. This illustrates what type of pattern mining methods can be used on a movement dataset.

6.2. Movement identification

By definition, moving objects are entities whose positions of geometric attributes change over time (Dodge, 2008). People always move in geographic space, this means that human movement is geo-referenced. When the start and end time of one movement is specified, its trajectory can be constructed by ordering several movements of one individual. These trajectories can be visualized and analysed.

In order to identify movement patterns, it is important to understand what types of patterns may exist in the data. Besides, there are many types of patterns and not everything is relevant for this project. Therefore, this section will organize various categories. This project aims to identify three different movement patterns: **1) Spatio-temporal movement patterns; 2) ordered co-location in space; 3) unordered co-location in space.**

Individual and group movement

Patterns can occur in individual movements or in movements of a larger group. Typical movements of individuals will be different from typical movements of a larger group. For analyzing movement in a larger area with more than 25.000 users, we are interested in typical movement at the larger aggregate level of crowds.

6.2.1. Spatio-temporal movement patterns

As described previous in this section, movement is from one location, or state, to another state, i.e. A to B. These movements can be analysed from movement data to detect the direct connectedness and flow between two locations in a time interval. Questions such as “where do people come from” and “how many people move between two locations” can be answered. Several patterns can be identified from this analysis. Firstly, the number of movements over time can be detected. This will provide insight in the behaviour of humans, e.g. when people go home or at what time people have lunch. Secondly, the flow and direction between two states, i.e. the analysis of the direction of the flows provides information on the symmetry of movement between two locations. For example, if movement 100 people move from A to B within a time interval and 100 people move from B to A in the same time interval, the movement pattern is perfectly symmetrical. Besides analysing movements between two states, consecutive movements of one individual can be used to identify movement patterns. These trajectories will be the basis for the next section to identify co-locations of several trajectories.

6.2.2. Co-location in space

When moving individuals share some locations in their trajectory, you can speak of co-location in space. According to Dodge (2008) there are three types of co-location in space: **1)** ordered co-location occurs when some locations are shared by multiple trajectories in the same order; **2)** unordered co-location when shared locations are attained in different orders; **3)** symmetrical co-location when the shared locations are in opposite order. This means that co-location in space, helps to identify movement patterns in the sense of frequently visited locations in one trajectory. For example buildings A, B, C can be visited in the same order by multiple trajectories, and the same buildings can be visited by multiple trajectories, but in different orders.

Ordered co-location in space can be analysed with the concept of sequences. A sequence is an ordered list of visited locations. Sequential pattern mining algorithm help to understand what order common locations are visited. In this report, trajectories of a sequence of locations are analysed to identify ordered co-location in space movement patterns. Unordered co-location in space analyses the same trajectories, but does not consider direction or order of the movement. This means that common locations visited together in one trajectory can be identified. In other words, the association between buildings is detected. A commonly used method to detect groups of objects in a list (i.e. a trajectory), an association rule mining algorithm is used. This report will use the concept of this algorithm to identify these groups of buildings that are frequently visited together.

7

Methodology

In this chapter the data mining methods used to retrieve movement patterns from the TU Delft eduroam Wi-Fi log data will be described in detail. Figure 7.1 gives an overview of the main workflow to derive movement patterns from the Wi-Fi log. First the raw Wi-Fi log is preprocessed to get states at two different spatial levels (building- and building-part level). A state is defined as a time interval during which a particular device is located in a certain area. An example of a state on building level is: device A is located at Library from 11:00 to 12:00. An example on building part level is: device A is located at canteen from 11:00 to 12:00. In the pre-processing phase the data is enriched with 'world' states, reduced by grouping states and cleaned by filtering out 5 minute states representing people that only pass by a building without actually entering it. The insertion of world states enables the detection of movement from and to the campus in the case of building level, and movement from and to the building in the case of building-part level. The states resulting from the pre-processing are used to retrieve movements at both spatial levels. A movement is defined by the change from one state to the next subsequent state, where the different states must be at a different locations. Furthermore, the building level states are used to retrieve trajectories for each device. A trajectory is defined as an ordered list of states. The trajectory thus stores the entire route or trajectory the particular person travelled. For the building-part level no trajectories are retrieved. For building-part level a graph is made for BK-city. In these graph the nodes represent the different building-part and the edge follow movement space, such as corridors and stairs. Using the shortest path in the graph, the route of the movements within BK-city can be visualized in more detail. For building level no graph is created and both trajectories and movements are visualized simply as a straight line. In addition to the maps at both the building and building-part spatial level, movement time series are created for both spatial levels. Together these maps and time series are used to identify different types of movement patterns.

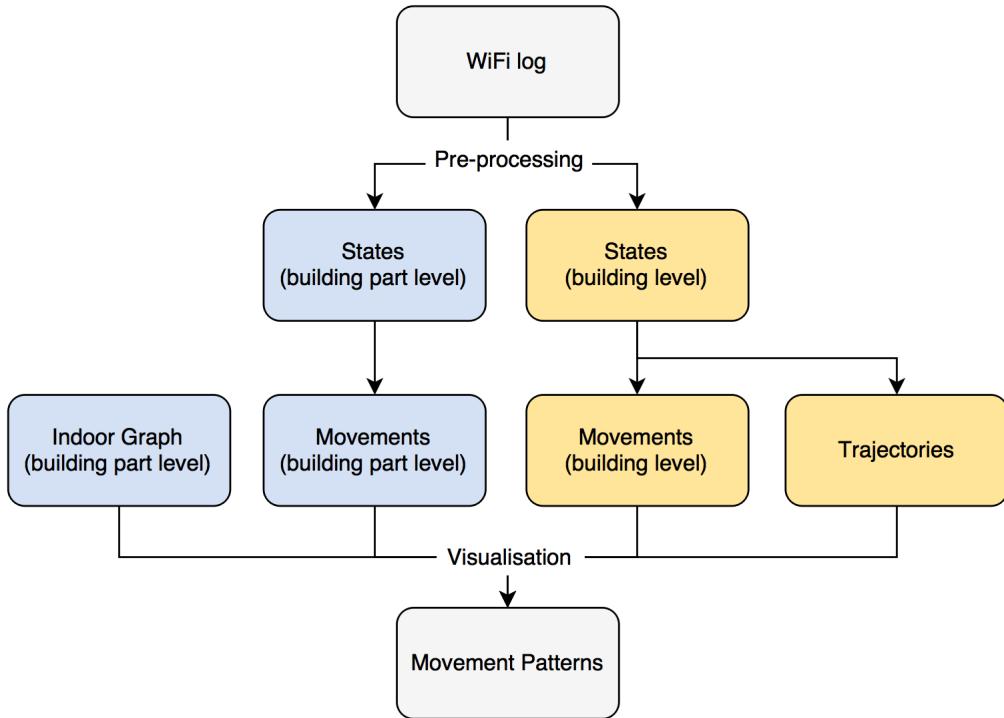


Figure 7.1: Grouping

In the following three chapter all these steps to derive different movement patterns will be described in more detail. First section 7.1 describes the various pre-processing steps to clean, reduce and enrich the raw data. Subsequently ?? addresses the creation and visualization of movement on building level, this includes movement between buildings and movement from and to the campus. Chapter 9 covers the creation and visualization of trajectories on building level. Finally chapter 10 discusses the creation of the graph and movements for BK-city (building-part level) and visualization of these movements. The latter three of these chapters also show several results. The focus of this project however is on developing methods to identify movement patterns. It should therefore be kept in mind that many more results concerning movement patterns on the TU Delft campus and inside BK-city could be retrieved using the described methods.

7.1. Preprocessing

Before movement patterns between buildings and building-parts can be retrieved, pre-processing of the raw data is required. In this chapter the different pre-processing steps will be described in detail. First section ... addresses the initial data filtering. Section ... concerns the grouping of records with the same mac address and location that are subsequent in time. Section ... describes the filling of the dataset with a 'world' location. This enables detection of movement from and to the campus on the building level, and movement from and to the building on building-part level. Finally section ... is about the filtering of records of people only passing by a building or building-part.

7.1.1. Initial filtering

Each record in the wifilog represents the scanning of a certain device at a certain time by a certain access point. In order to detect the movement patterns of these devices between buildings it should be known for each access point in which building it is located. The apname field in the wifilog table includes the building id in which building each scanner is located. However for some access points the apname is given in a different format and as a result their location is unknown. These apnames have in common that they don't contain the '-' character which is present in all the other apnames. As a result the apnames of which the location is not known can simply be filtered out by checking if a '-' is present in the apname. Two other special cases are present in which the building id in the apname is 1 or 102. These id's corresponds to the legermuseum and ... respectively. As the legermuseum is not located on the campus it was decided to omit this building. For

building-part level the general filtering is much easier, as only the records of a single building are of interest. Therefore, the records that don't have the building id of the particular building can simply be filtered out.

7.1.2. Grouping of states

In order to reduce the data and to be able to filter out records of people only passing by a building, the data needs to be grouped. The overall goal is to identify movement patterns between different buildings or building-parts. As a result records of subsequent states of the same device in the same building or building-part can be grouped together into one single record. Namely, if two subsequent states are at the same location they don't represent a movement and can thus be grouped. When looking at building level, the mobile of someone who studies the whole day at architecture might have 20 records (states) in the database for that day. This can be reduced to one record (state) that contains the time the device arrived at Architecture and left again. For building-part level the same applies for someone that has multiple subsequent states in the same building-part. To determine whether two records are subsequent in time, and therefore should be grouped together, a threshold for the time gap between two records needs to be defined. It was decided to set the gap threshold for grouping states to 1 hour. The reasoning behind this is that someone who is not scanned for a period of more than 1 hour has likely left the building. However, if someone is away for less than an hour it is more likely that that person was just smoking or lunching outside or just disconnected from the system for a while. Figure 7.2 gives an example of how the records are grouped on building level for one device for one single day. It should be noted that two states remain at faculty A as the gap between 12:30 and 13:45 is bigger than an hour. In the other other cases the gap is smaller than an hour and the records are grouped together. Only one record is present at faculty B so this record can not be grouped. The grouped records still contain all the information that is required to know that the person moved from faculty A to B to C during the day.

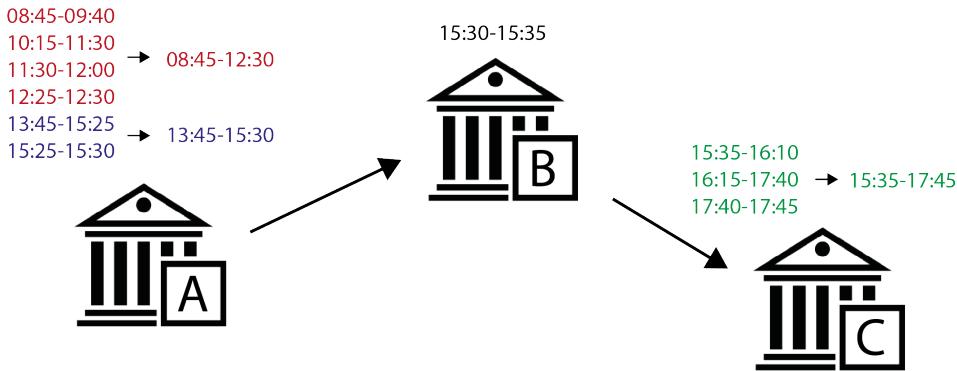


Figure 7.2: Grouping

7.1.3. Adding world state

Because the dataset contains all records of when certain devices are scanned, it also implicitly stores information on when the device is not located at the campus. These time gaps in which a particular device is not scanned at the campus give information on when the corresponding person is not at the campus. This information is valuable for detecting movement patterns from and to the campus in addition to the movement between buildings at the campus. Considering the fact that many students only visit one faculty each day. It becomes especially clear, that the movement from and to the campus plays an important role in the overall movement pattern of a person. In order to be able to directly derive movement from and to the campus from the dataset, the time gaps present in the data should be stored explicitly. Therefore each time gap larger than an hour is filled with a 'outside campus' or 'world' record. The word 'world' is used to indicate that the device could be located at any place in the world outside the campus during the time spans that it is not scanned at the campus. It should be noted that the reason that a device is not connected to one of the access points could also be that the device is simply switched off, in this case however the assumption is made that the device moves off campus. The begin and end time of a world record is defined by the end of the previous record and the start of the next record in time. In case there is no previous or next record the boundaries are defined by the starting time of the whole dataset and the current time. Figure 7.3 visualizes the explicit storing of world

states that fill time gaps during which a device is not recorded on campus. It can be seen that three 'world' states are added in the example. First during the start of the day before the person goes to faculty A, second during the lunch break, and finally in the end of the day starting from the moment when the person leaves faculty C. Storing these world states explicitly enriches the data as much more movement can be defined. The grouping of records described in subsection 7.1.2 and adding of a world state are complementary. If the gap between two states is smaller than an hour they are grouped if the gap is bigger than an hour a world state is inserted. For the building-part level the insertion of world states works exactly the same. In this case however the world represent the outside building area instead of outside campus area.

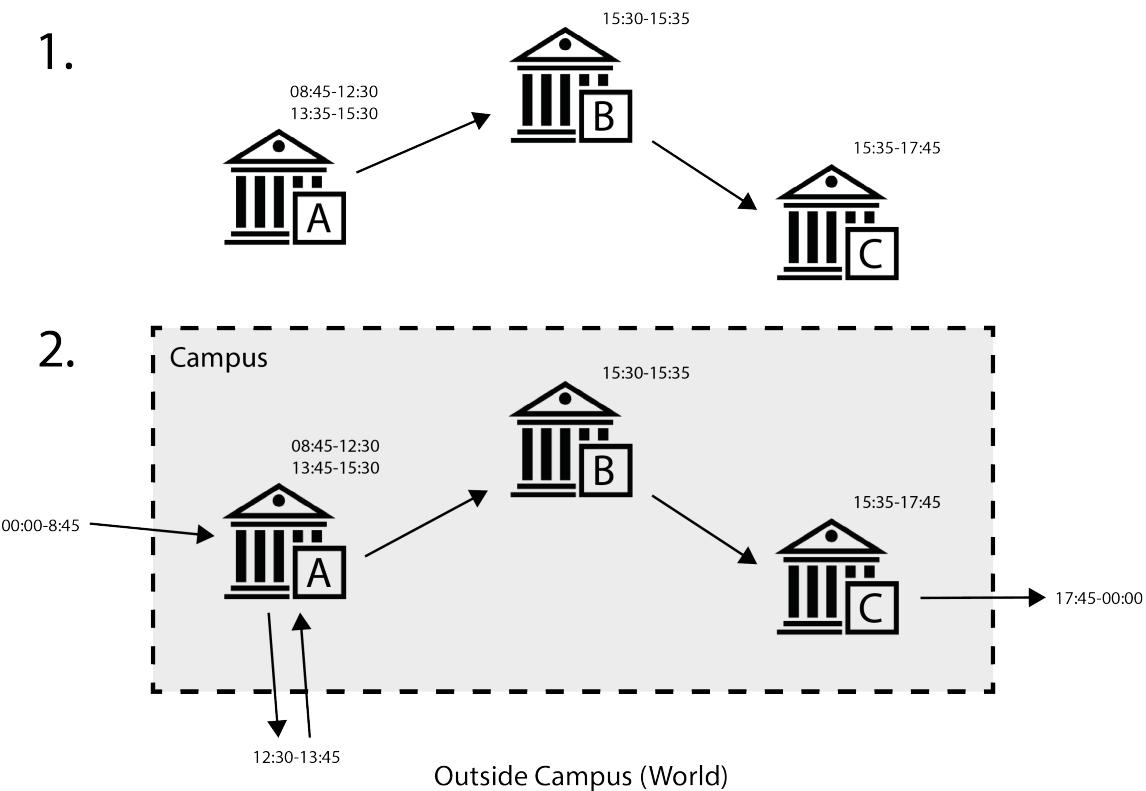


Figure 7.3: Adding World

7.1.4. Passing by

For the detection of movement patterns between buildings, records of people that only pass by a building without actually visiting it should be excluded. The reason for this is that records of people only passing by a building could result in misinterpretation of the movement patterns. This is illustrated by the example in Figure 7.4. In this case faculty B is located on the route from faculty A to faculty C. Therefore it is likely that people moving from faculty A to faculty C are picked up by a scanner located at faculty B. The eduroam system records all devices at intervals of approximately 5 minutes as explained in subsection 5.6.1. Such a recording by the eduroam system can happen during the short time period that the device, which is on its way from faculty A to C, is connected at faculty B. This will result in records of approximately 5 minutes at faculty B, whilst the person has not been inside faculty B. As the movement is the change between two states, the movement to faculty C will originate from faculty B. Someone that is not aware of the 'passing by' problem might conclude that people from faculty B often go to faculty C. In reality however, people from faculty A go often to faculty C. By filtering out the records of people only passing by buildings the correct movement can be visualized (see Figure 7.4 bottom). It should be noted that filtering out 'passing by' records can only be done after the grouping process. The reason for this is that 5-minute records that would individually be classified as someone passing by might be grouped together into one record with a longer duration. After grouping the combined record is not classified as someone who passes by anymore. Furthermore it should be noted that the filtering of 'passing by' records occurs after filling the data with 'world' states. The reason for this is that a passing by event does mean that the device was located on the campus. The world records

are meant to represent the time the device is not on the campus. Filtering passing by events works exactly the same for the building-part level. If a person only passes by a particular building-part without staying in it, it is filtered out. For building-parts this filtering is especially important as the route from one building-part to another often leads through several other building-parts.

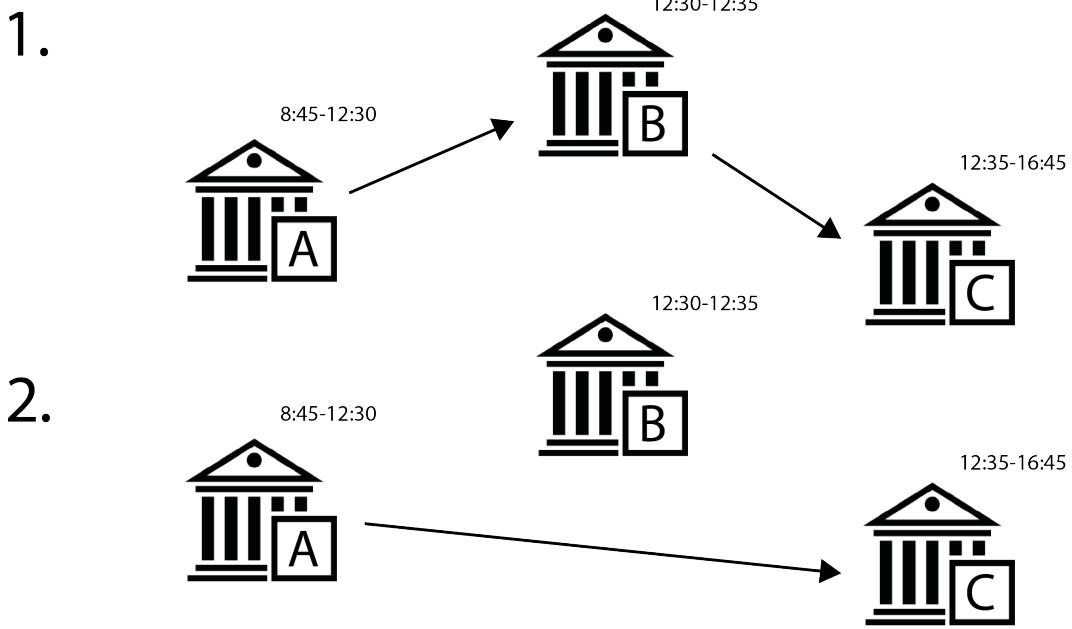


Figure 7.4: Passing by

7.1.5. Implementation

The filling (world), grouping and filtering (passing by) steps described above are implemented in an integrated way. The Pseudocode for the implementation for building level is shown in Figure 7.5, for building-part level the implementation is exactly the same only grouping is done for building-parts instead of buildings. As can be seen in the code there is communication with the database at several points. The table from which the records are retrieved for each mac address is already processed as described in the general filtering section. Furthermore the format of the table is slightly different compared to the initial wifilog. The session duration is exchanged for an end time column which is derived by adding the session duration to the astime (start time of a record).

```

macs = get distinct macs from db
create new empty table with 4 columns(mac, building, start, end)
min_time = minimum time in entire db
max_time = current time
for mac in macs:
    records = get all records for mac from db
    cur_rec = first record from records
    insert world at start (mac, world, min_time, cur_rec[start])           # fill
    for next_rec in records[1:-1]:
        gap = next_rec[start] - cur_rec[end]
        if gap > hour:
            insert world (mac, 'world', cur_rec[end],next_rec[start])       # fill
        if gap < 15 minutes and cur_rec[building] == next_rec[building]:
            cur_rec = (mac,cur_rec[building],cur_rec[start],next_rec[end])   # group
        elif cur_rec[end]-cur_rec[start] > 6 minutes:                      # filter passing by
            insert cur_rec
            cur_rec = next_rec
        if cur_rec[i_end]-cur_rec[i_start] > 6 minutes:                      # filter passing by
            insert cur_rec
    insert world at end (mac, world, cur_rec[end],max_time)                  # fill

```

Figure 7.5: Pseudocode preprocessing

Figure 7.6 shows an example of the records of one device over a time span of one day during the different pre-processing steps. From the raw data it can be seen that this person spends most of the day in building B. The person is scanned once at building A before he arrives in the morning and after what is likely to be his lunch break. The last two hours the person is scanned in building C. After filling three world records are added, at the beginning of the day, during the lunch break, and at the end of the day. The grouped records show that the subsequent scans in building B and C are grouped together. Finally the scans at building A are removed from the dataset as they are likely to indicate passing by events.

Raw			Filled		
Bld.	Start	End	Bld.	Start	End
A	09:30:00	09:35:07	W	00:00:00	09:30:00
B	09:35:07	09:40:07	A	09:30:00	09:35:07
B	09:50:28	10:41:21	B	09:35:07	09:40:07
B	10:41:21	12:08:40	B	09:50:28	10:41:21
B	12:08:40	12:13:51	B	10:41:21	12:08:40
A	13:30:03	13:35:12	B	12:08:40	12:13:51
B	13:35:12	13:40:16	W	12:13:51	13:30:03
B	13:40:16	15:34:22	A	13:30:03	13:35:12
B	15:34:22	15:39:26	B	13:35:12	13:40:16
B	15:44:34	15:49:34	B	13:40:16	15:34:22
C	15:59:47	18:06:54	B	15:34:22	15:39:26
C	18:06:54	18:11:54	B	15:44:34	15:49:34

Grouped			Filtered (Passing by)		
Bld.	Start	End	Bld.	Start	End
W	00:00:00	09:30:00	W	00:00:00	09:30:00
A	09:30:00	09:35:07	B	09:35:07	12:13:51
B	09:35:07	12:13:51	W	12:13:51	13:30:03
W	12:13:51	13:30:03	A	13:30:03	13:35:12
A	13:30:03	13:35:12	B	13:35:12	15:49:34
B	13:35:12	15:49:34	C	15:59:47	18:11:54
C	15:59:47	18:11:54	W	18:11:54	00:00:00
W	18:11:54	00:00:00			

Figure 7.6: Preprocessing

7.1.6. Apname vs maploc

The data in the table 'wifilog' contains information about the location of the Access Point (AP) in two columns. The first one is the column 'apname', which is a string with the symbolic name of the AP, for example 'A-08-G-010'. The two numbers in the second part of the string, in this case '08', represent the building number. This building number can be linked to a location in the world. The second column which contains information about location, is the column 'maploc'. This column also contains strings, which look as follows:

System Campus > [buildingid] > [specific location]. An example of such a string is 'System Campus > 21-BTUD > 1e verdieping'. In such a string, the middle part can be linked to a building, so to a real-world location. But there are some other values for maploc, which can less clearly be linked to a real-world location. Such a value is 'Root Area', it is unclear what this value means and it contains no information about a building or area it might be in. This makes it impossible to link it to a location in the world. Then there is the value 'Unknown', a value that indicates that there was no name attached to the Access Point that user was connected to. Again in this case, it is impossible to link this value to a real-world location.

As both 'Root Area' and 'Unknown' are in the minority of records, they could be left out of the queries. But for some records, the column 'apname' did provide information about the location, while the 'maploc'

column value was 'Root Area'. In most of these cases however, the building number, the second part of the string, was a number of length three. But there are no buildings on the TU Delft campus with a building number that high. When consulting Wilko Quack about this, he explained that these building numbers had an arbitrary 1 in front of the building number. So 'A-134-A-001' was not building 134, but building 34, which was an actual building number on the campus. This would mean that using the column 'apname' for getting the building number would mean a higher number of results and therefore a more realistic visualization of the movements.

Taking the substring of that column and linking it to a building with an actual location is done in two steps. First the whole string is retrieved and with a function in Python the substring is derived. Subsequently, the building id that is the result of this function can be linked to a table in the database which has for every building five columns: buildingid, name, point (as geometry), x (longitude), y (latitude) (see in ??).

7.1.7. Static and mobile devices

In order to identify the movement patterns and know what entrances and exits are most frequently used even better, we aim to identify dynamic and static devices. In our first approach, we will look at the number of different access points the device is scanned by in time. The distinction between static and dynamic devices is important, because the behaviour, in terms of Wi-Fi tracking, is significantly different. For instance, a static device, such as a laptop, connects with the Wi-Fi network at different moments compared to a dynamic device, such as a mobile phone. The difference will be explained more in detail using the image below.

Assume a person that carries a static device (laptop) and a dynamic device (mobile phone) enters a building. While being on his way to the destination, the person does not make use of the laptop, thus the laptop is not connected to the Wi-Fi network. On the other hand, the Wi-Fi of the mobile phone is turned on all the time, and connects at the moment the device is on range of the first access point. On the way the mobile phone is scanned by Access Point(AP) 1, 2 and 3. The person connects to the Wi-Fi network with the laptop at the moment it arrives in the room, of which the Wi-Fi is covered by AP 3. This access point scans the laptop for first time after entering the building. The static laptop is distorting the result, due the fact that in this case the entrance access point for the laptop would be AP 3. In order to achieve a more reliable result, the aim is to filter out the static devices.

To identify the static and dynamic devices, we analyze the behaviour of each device. The first approach focuses on the number of (distinct) access points and the session duration. We assume to find differences between them (Table 7.1).

	Session duration	Nr.of access points
Static	long	low
Dynamic	short	high

Table 7.1: Difference between static and dynamic devices

We expect that the relation between the distinct access points and the (summed) session duration, called ratio, is going to be useful in making the distinction between static and dynamic devices (Equation 7.1).

$$\text{Ratio} = \text{distinctaccesspoint} / \text{summedsessionduration} \quad (7.1)$$

In this, a small ratio indicates the device is dynamic and a large ratio indicates the device is static. The result shows that the number of devices decreases over ratio(Figure 7.7).

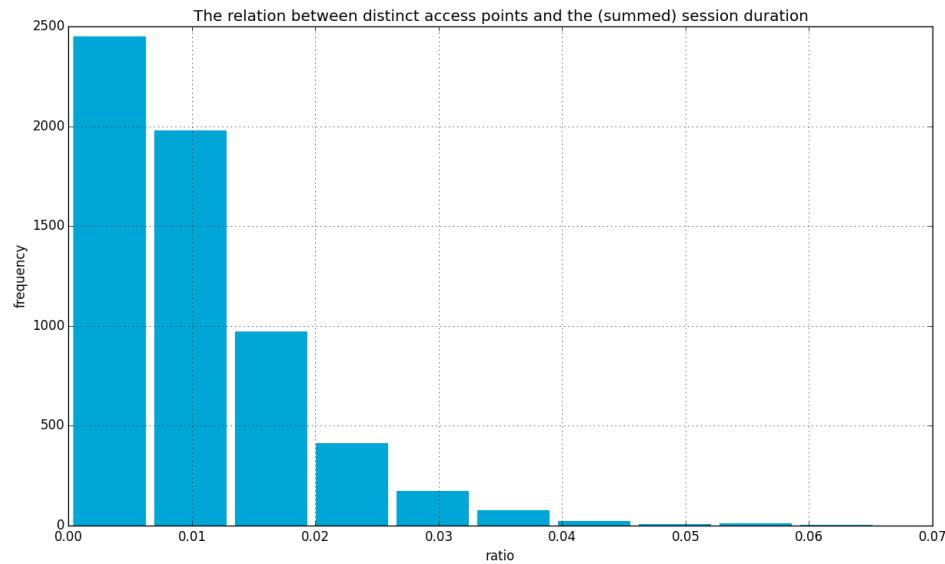
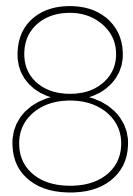


Figure 7.7: The relation indicating frequency of a radio

Because the frequency decreases gradually, there is a fuzzy boundary that separates the static from dynamic devices. Therefore it is not (yet) possible to filter out the static devices for further analysis. In order to improve this, the plan is to use the exact number of access points that scanned the device instead of the distinct access points. Also, a closer look will be taken at the session duration, since dynamic devices will have session duration of approximately 5 minutes much more often.

7.2. States to movements



Spatio-temporal movement patterns

8.1. Introduction

8.2. Methods

8.3. Results

8.3.1. All movement

8.3.2. Mobile vs static

8.3.3. Week vs weekend

8.3.4. From and to

Campus Library and BK Aula during lunchtime Cantine and bouwpub

9

Trajectory patterns

9.1. Introduction

This GSP attempts to identify people's movement patterns from anonymized wifi logs. chapter 8 described movement patterns including spatial and temporal aspects of single movements of a crowd of people. Another way of looking at movements, is by tracking individual movement for a longer time interval. A large set of individual trajectories can be used for the identification of typical movements among users of the campus. The method uses concepts from sequential pattern mining.

This chapter presents a method for identifying movement patterns using individual trajectories. As described in chapter 6, if moving individuals share some locations in their trajectory, you can speak of co-location in space. When the order of the shared locations are similar for multiple trajectories, you can speak of typical movement. This concept is explored for the identification of movement patterns, and thus the usage of the campus. This approach can answer different questions than looking at single movements, as is done in chapter 8. For example, 'how many places the user frequently visits', 'at what order the user visits places', 'how often a trajectory happens', 'how many places contained in a frequent trajectory'.

First, this chapter will describe the problem description, including the extraction of locations of a user, the mining of individual trajectories from an anonymized Wi-Fi scan list, and finally the mining of movement patterns from a set of trajectories using the PrefixSpan algorithm.

9.2. Problem description

9.2.1. Location extraction

The data provided by the eduroam network enables a detailed view of people's movement on campus. The large coverage of the eduroam network allows to track users for a large part of the day when they enter the campus. However, the observation space is limited to the extent of the size of the campus, making it not possible to track people outside the eduroam network. A second disadvantage is the spatial resolution of the positioning method. The range a mobile device can be connected to an AP, influences the accuracy of the estimated location of a mobile device. For indoor environments of the TU Delft campus, this is just a few tens of meters wide. This resolution allows tracking movement at a building level by re-locating mobile devices to the closest AP. Data between two re-locations is not available. Therefore, an individual's trajectory is depicted by connecting the re-locations as a sequence of APs. These individual trajectories are used to identify patterns.

A location represents a geographic position where a user stays, i.e. a user is in state. For identifying movement patterns from Wi-Fi monitoring, we are interested in movement between two locations where an individual stays for a longer time period. Such a location, or stay place, can be detected when a user is connected to the same AP for a longer time. To detect buildings as a location (i.e. contains multiple APs), two consecutive Wi-Fi scans must contain APs of the same building. With a data collecting interval of 5 minutes, it means that people will be filtered out if their stay duration is less than 10 minutes. Based on this assumption, people with a shorter stay duration are considered passing by, as explained in section 7.1.

9.2.2. Individual trajectory

An individual's trajectory is constructed as a sequence of locations in order of the scan time. Start and end time of a trajectory can be specified with a time interval. Two consecutive scans from the Wi-Fi log are considered in the same trajectory if and only if $t_{s2} - t_{e1} < T_{split}$, where T_{split} is the splitting threshold. The splitting threshold is important when dealing with people, who are not observed for a long duration of time, i.e. people moving home. For example, if a student leaves the campus at the end of the day, and returns the next morning, separate trajectories should be created. Because, T_{split} is larger than the threshold for identifying '*world*' (see section 7.1), the trajectory will always start and end with '*world*'. If p is a location, then a trajectory can be written as:

$$p_1 \rightarrow p_2 \rightarrow p_3 \rightarrow \dots \rightarrow p_n$$

Given a time interval, there is a set of individual trajectories $S = \{t_1, t_2, t_3, \dots, t_n\}$ where each t_i is the trajectory.

9.2.3. Trajectory Pattern

From a set S of trajectories, different patterns can be identified using sequential pattern mining algorithms. Frequency of a trajectory by all users of the campus can be detected. This can be represented as a trajectory T with a support s . Support means how many times the same sequence, or sub-sequence, is shared in the set of trajectories. This gives valuable information on the order common buildings are used and what order of buildings occurs the most. Using a minimum support threshold, sequential mining returns all movement patterns that satisfy $n > 2$ and support $T > S_{min}$. Furthermore, the length of common trajectories can be discovered. This allows for identification of movement patterns of a specific length n . Also, when location is not considered, but only the length of a trajectory, the mobility pattern of an individual can be described in terms of how many times he/she re-locates. Figure 9.1 illustrates a trajectory pattern of length 3, and has a support of 3.

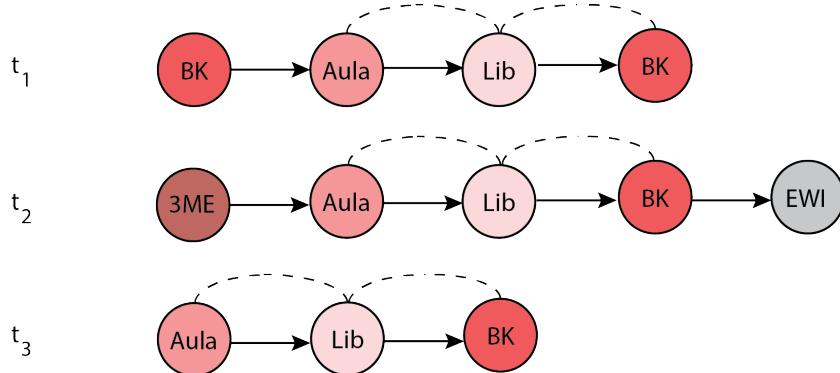


Figure 9.1: sample of individual trajectories

For this study, a trajectory pattern is a sequence of states with $n > 2$ and support $> S_{min}$. We are only considering trajectory patterns with $n > 2$, because chapter 8 already looked at two consecutive states.

There exists many developed sequential pattern mining algorithms. For this study PrefixSpan Pei et al. 2004 is used to identify common shared trajectories or sub-trajectories. This sequential pattern mining algorithm can find re-occurring sequences or sub-sequences from a set of trajectories. For every common sequence, a support value is computed.

9.3. Implementation

For this analysis the same data is used as for the single movement analysis. As described in section 7.1, from the raw Wi-Fi log states are extracted. More than 2.8 million states are identified from the dataset. This information is stored including a unique mac address, a number representing a building, the start time of the state and the end time of the state. The states are used to construct individual sequences ordered by date and time. The T_{split} is used to create separate trajectories for different days, for each individual. For this study, a new trajectory is created when there has not been a connection for 5.5 hours, i.e. a state of outside campus ('*world*') > 5.5 hours. This threshold is suitable for identifying people moving home at the end of the day and coming back the next morning. After splitting the sequences, over 950.000 trajectories are created,

with temporal granularity of one day. Every trajectory starts and ends with 'world', i.e. people start and end their trajectory outside the campus. A sample of a constructed trajectory can be seen in Figure 9.2

```

World → Bk → World → Aula → World
World → 3ME → Lib → Aula → Lib → World
World → EWI → World

```

Figure 9.2: sample of individual trajectories

Based on the created trajectories, trajectory patterns with a support value are detected by applying the PrefixSpan algorithm. Figure 9.3 shows an example of the detection of patterns with a support value given by the sequential pattern algorithm. Logically, the pattern with the highest support is a length-1 sequence. The longer patterns get, the lower the support will be.

0 23 0	([0, 21], 4)
0 23 0	([21, 0], 4)
0 32 0	([21, 21], 4)
0 22 20 0	([0, 0, 21], 4)
0 22 0	([0, 21, 0], 4)
0 21 0 21 0	→ ([0, 21, 21], 4)
0 21 0 21 36 0	([21, 0, 21], 4)
0 21 0 21 0	([0, 21, 0, 21], 4)
0 36 0 36 0	([0, 0], 5)
0 21 0 21 0	([0], 10)

Figure 9.3: sample of individual trajectories

9.4. Results

10

Indoor spatio-temporal movement patterns

10.1. Introduction

As described in the first part of this report, Wi-Fi tracking data can be used to identify movement between buildings. Given that indoor areas are usually better covered with Wi-Fi access points than outdoor areas, it is natural to also look at movement inside buildings. The following section describes our method of identifying and visualizing indoor movement in the Faculty of Architecture of TU Delft.

The process of indoor movement analysis is conducted along the steps below, thus the section also follows this structure:

1. Delineate building parts based on the layout of access points and the division of the building (e.g. department, canteen, building wing), and group the access point into building parts.
2. Identify movements in the data between building parts.
3. Create a route network that connects the building parts and is constrained on the corridors of the building.
4. Assign the movements to the route network.
5. Visualize the movement along the indoor network.

10.2. Theory / methods

After identifying movement between different buildings, the next level is to do so between different parts inside a building. These parts represent functional or spatial divisions inside a building, e.g. departments, community areas, building wings and are referred to as *building part*.

A prerequisite of the method is to know the at least room level location of the access points in the respective building. At the time when the project was carried out, the detailed access point locations were available only for the Faculty of Architecture. Thus the focus on this particular building.

As opposed to outdoor pedestrian movement which is not necessarily constrained on a fixed network, indoor movement is constrained by the layout of the respective building. The building parts of the Faculty of Architecture can be represented by its underlying graph, having the building parts as nodes and the corridors as edges Figure 10.1. Then indoor movement is necessarily constrained on this underlying graph.

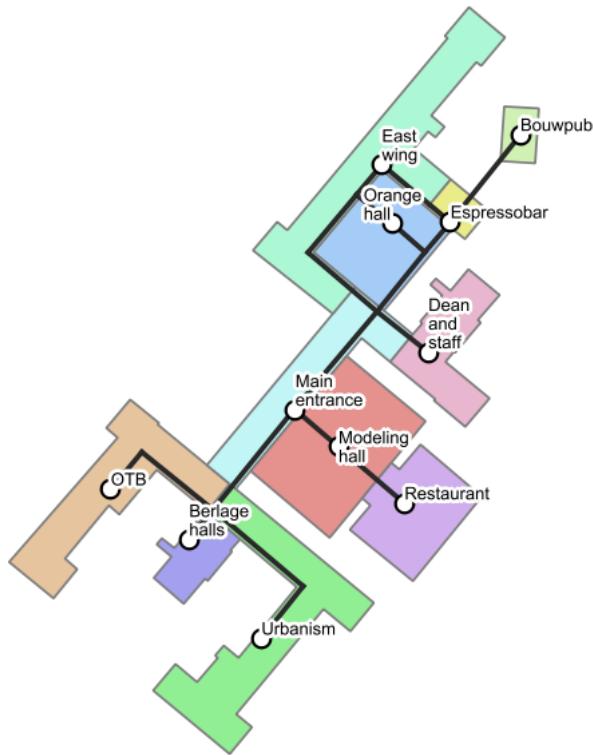


Figure 10.1: Building parts on the ground floor of the Faculty of Architecture and its underlying graph.

The Wi-Fi system of the TU Delft campus has a five minute scan interval, which is too coarse to catch detailed movement indoor. As five minutes is sufficient to reach any two locations in the building taking any route. Therefore not the movement trajectory itself is identified from the data, but the fact of relocation from origin to destination. Then the path of the movement can also be identified by analysing the layout of the building. For example if a person stayed at the Restaurant, then soon after he stayed at the Orange hall, he necessarily had to traverse the corridors in-between these two locations. Our method is based on this assumption.

Due to the building layout, in most of the cases there is only one possible direct route between two building parts. However, in case of multiple route options, the exact route of a movement is assumed to be the shortest route between origin and destination.

10.3. Implementation

The identification and visualization of indoor movement ins a procedure that requires various tools and steps. While some steps can be automated, others need to be done manually. The detailed description of these steps follows.

10.3.1. Delineation of building parts

There are two factors that define what is considered a building part, the layout of the building and the layout of the access points. The layout of the building defines the functional divisions, e.g. departments or common areas. Additionally, it is necessary to have at least one access point in each of these divisions, or preferably more access points equally distributed in the division. Considering the signal range of an access point, it is not desirable to have access points close to the boarder of two neighbouring divisions, as in that case the user could be falsely located in the neighbour division if he is picked up by the respective access point. The combination of a functional division and the access points within define a building part.

In case of the Faculty of Architecture Figure 10.2 displays the provided access point map and the manually overlaid functional divisions, thus defining the building parts.

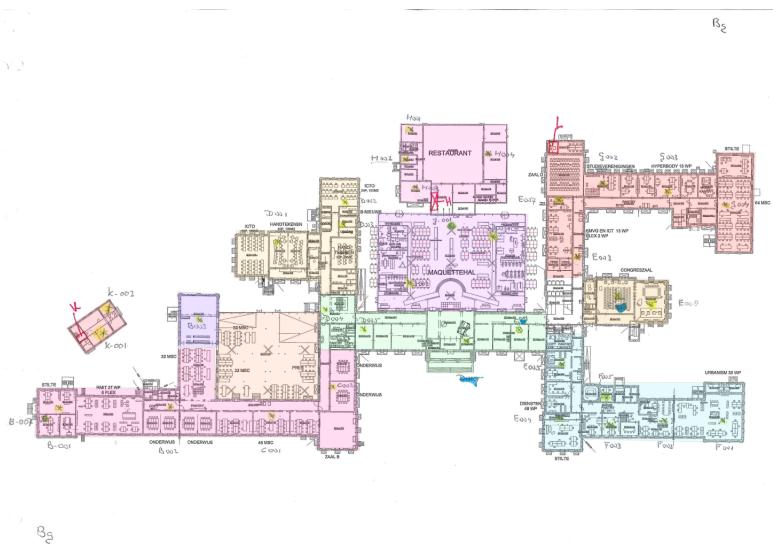


Figure 10.2: Access point map where yellow dots mark the access points, and the functional divisions (coloured areas) on the ground floor at the Faculty of Architecture.

10.3.2. Movement between building parts

The method how indoor movements are identified is described in REFER TO THE CORRECT (SUB)SECTION

10.3.3. Indoor route network

The route network of the Faculty of Architecture where nodes represent building parts and edges represent corridors was drawn manually in QGIS, following the floor plan of the building. However, the resulting *spaghetti network* does not contain the topological relations that are required to calculate a shortest route. Therefore the topological relationships were created with the PostGIS extension *pgRouting*. Using a database-based solution for storing the data, creating topology and calculate shortest routes allowed us to easily match the movements, which were calculated in the database, to the route network.

10.3.4. Mapping traffic to the route network

In the *movements table* every record represent a single move of a person from origin to destination. In order to display these movements, identical moves that have the same origin-destination pair are aggregated, resulting in a table of unique origin-destination pairs with the amount of related moves Table 10.1.

Origin	Destination	Count
OTB	Restaurant	126
Main entrance	Espressobar	543

Table 10.1: Aggregated moves between building parts

Then the shortest route between each origin-destination pair is calculated and the movement counts are added to each edge that is traversed in the network. Thus if the shortest route of two distinct movements share edges, the movement count is summed up on the common edges, resulting in the traffic load of a given edge Table 10.2.

Edge ID	Traffic	Line width
45	6151	1.10
46	1994	0.64

Table 10.2: Traffic load on the indoor route network

10.3.5. Visualization of the movement

The visualization method, as well as the route network, is two-dimensional. However, three-dimensionality is imitated by using an *exploded view* common in architectural visualizations, that shifts overlapping elements (e.g. floors) by a certain angle.

In this graphic the *nodes* that represent the building parts are the approximate centroids of the polygonal area of the building part. The nodes were manually adjusted to better match the route network.

The route network is represented with straight lines, where the *line width* is proportional to the traffic load of a given edge. However, line widths cannot be compared across graphics, as in order to facilitate consistent scale the line width variable is normalized to the range of 0.5-5 units, regardless of traffic load. The range of 0.5-5 units is chosen to provide a visually appealing and clear graphic. *Colours* mark the four separate floors and the staircases (grey) in the building Figure 10.4.

10.4. Results

Considering the movement from any origin to any destination at the Faculty of Architecture, throughout the whole measured period, our results clearly indicate the peak hours in the morning, before and after lunch Figure 10.3.




Figure 10.3: Total movement at the BK

Furthermore, if the same data is visualized using the previously described method, we can observe the occupation of corridors in the building. The advantage of this method that it provides insight into the usage of those spaces where data is not directly available. See the limitations of the *eduroam* system to track detailed indoor movement in section 10.2.




Figure 10.4: Occupation of corridors at the BK

Additionally, we analysed movements of mobile and static devices, weekdays and weekends, to the Bouwpub and to the Restaurant at the Faculty. These results are displayed in chapter 14.

11

Conclusions

First of all it can be concluded from the preliminary results that the Wi-Fi network data is suitable, at least to some extent, for retrieving movement patterns of people. Expected patterns such as a movement peak between building during lunch time, and a morning and afternoon peak of people entering and leaving the campus can be clearly distinguished in the data. Similarly aggregated movement on the map shows the expected result that Aula-Library is the most frequently travelled path. More specific patterns between particular buildings and/or during certain time intervals can easily be derived due to the automated workflow. An example of such a specific pattern is that people moving to the aula most often origin from the faculty of Applied Sciences. Furthermore it can be concluded that Aerospace Engineering and to some extent Architecture are rather isolated compared to the other faculties on the campus. Especially when interpreting the result of movement from and to the campus, it should be taken into account that static devices (mainly laptops) are not filtered yet. Disconnecting a laptop for over an hour will currently still be interpreted as a movement away from the campus and back.

12

Recommendations

12.1. Entrances

12.1.1. Introduction

This section will describe the work that is done to find out what, when and how frequent entrances of the Faculty of Architecture are used. This is an interesting and challenging use case at the same time. The Faculty of Architecture is a building having multiple entrances; five to be precise. Knowing what, when and how frequent these entrances are used, will give insight into the use of a building, the spatial context and the relation between these two.

12.1.2. Methodology

In order to find what entrance someone uses to enter or exit a building, we will look in the part of a sequence in which the device is recorded by an AP in a building and subsequently recorded by an AP in another building. More specific, we will look at what (first or last) AP is used in a movement from one to another building. For this two different approaches can be distinguished. The first approach does not take in account the devices that might get recorded when passing by the building. In the second approach we will make use of the pre-processed data which excluded the passing by events.

12.1.3. Hypothesis

Our hypothesis is that finding clear answers to the question whether it is possible to identify what entrances are most frequently used, is going to be hard. Firstly, because the existing layout of APs is not designed for the purpose of tracking people. For this reason there is not always an AP located near an entrance. Secondly, because the logging frequency of the system is a little more than 5 minutes. Ideally the system records the connected device at the very first AP it connects with. The chance the device is recorded at the moment it is connected with the very first access point is small. However we still expect to see some results. Although the time interval in which the system logs the connected devices is relatively large, an AP located near an entrance would still pop up as one of the most frequently used AP as first connection (assuming people disseminate over the building after entering).

12.1.4. First approach: including passing by events

The first approach makes use of the raw wifilog data, by finding the part in a sequence in which a device is recorded by an AP in a building and is subsequently recorded in another building. The states in which a device is scanned once are not filter out. These single records imply that a device only passed by the building, and thus was not located in the building.

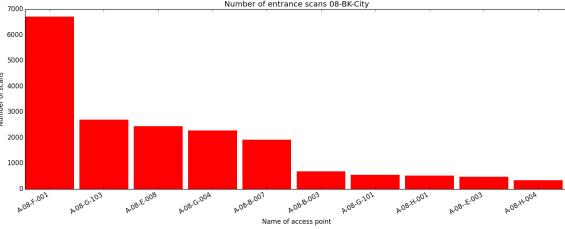


Figure 12.1: Most frequently recorded APs in a movement to the Faculty of Architecture

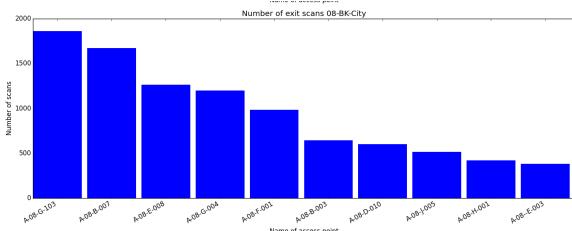


Figure 12.2: Most frequently recorded APs in a movement from the Faculty of Architecture

The floor plans of the Faculty of Architecture, enriched with the location of APs, are used to locate the most frequently used APs on the map (see ??). The result is interesting, since most APs are not located near an entrance but are located at one of the corners of the building. Most of them are located at the western part of the building. Knowing that lots of people are passing in the street next to this part of the building, we can conclude the result of this analysis is distorted due not filtering out the devices that are recorded when passing by the building.




Figure 12.3: The location of the most frequently used APs that are used to record the first and/or last connection of a device in the Faculty of Architecture

12.1.5. Second approach: excluding passing by events

Table x shows the individual states as a result of the pre-processing (see chapter pre-processing). The records represent the states for each mac, including the first and last recorded AP (ap_start, ap_end).

mac	building	ts	te	ap_start	ap_end
000c+YfkIi..	0	30-3-2016 23:34	6-4-2016 22:39	NULL	NULL
000c+YfkIi..	21	6-4-2016 22:39	6-4-2016 23:30	A-21-0-005	A-21-0-045
000c+YfkIi..	0	6-4-2016 23:40	10-4-2016 19:53	NULL	NULL
000c+YfkIi..	0	10-4-2016 20:03	10-4-2016 21:13	NULL	NULL
000c+YfkIi..	21	10-4-2016 21:13	10-4-2016 21:34	A-21-0-046	A-21-0-046
000c+YfkIi..	21	10-4-2016 22:04	10-4-2016 22:19	A-21-0-045	A-21-0-046
000c+YfkIi..	0	10-4-2016 22:19	10-4-2016 23:14	NULL	NULL
000c+YfkIi..	0	10-4-2016 23:24	11-4-2016 12:27	NULL	NULL
000c+YfkIi..	21	11-4-2016 12:27	11-4-2016 13:25	A-21-0-043	A-21-0-043
000c+YfkIi..	20	11-4-2016 13:25	11-4-2016 13:56	A-20-0-008	A-20-0-045

Table 12.1: Individual states as a result of the pre-processing

The table also includes 'world' (in Table 12.1 represented by NULL) which implies the device is not located on the campus. A simple SQL query is used for plotting the most frequently used first and last recorded APs in a stay (Figure 12.5)

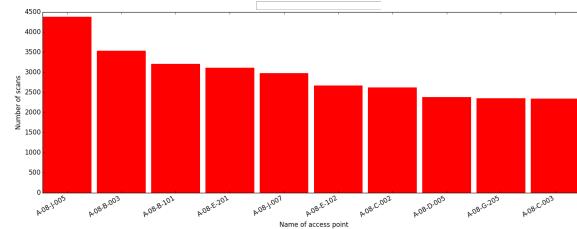


Figure 12.4: Most frequently recorded APs in a movement to the Faculty of Architecture

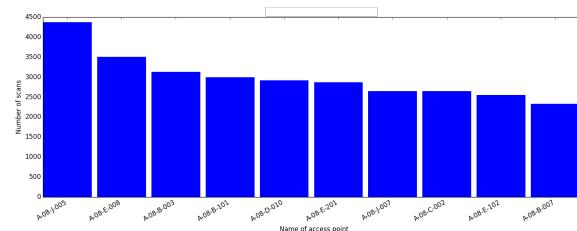


Figure 12.5: Most frequently recorded APs in a movement from the Faculty of Architecture

The most frequently used access point, A-08-J-005, is located high up in the modelling hall and thus not near an entrance (see Figure 12.6). Although this location is different than expected there might be a reason for it. The access point is placed in an open space in which no objects could seriously block the Wi-Fi signal.

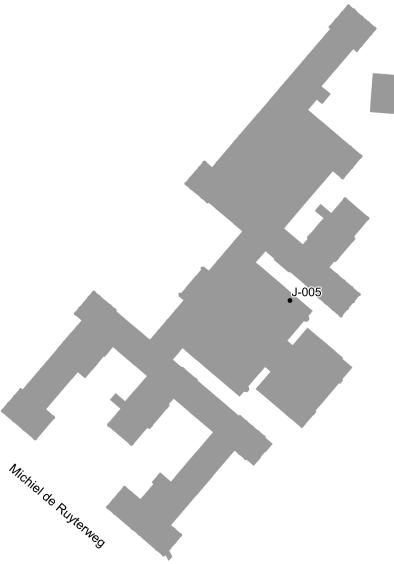


Figure 12.6: The location of the most frequently used APs that are used to record the first and/or last connection of a device in the Faculty of Architecture

In order to know with what APs a device connects when entering a building, some experiments are conducted. By looking at the MAC address of the access point the device connects with, it would be possible to identify the location of that AP. This experiment is conducted for entering the Faculty of Architecture via the East, West and main entrance. Table 12.2 shows the results of the experiment.

entrance	MAC address	apname	maploc
east entrance	00-15-C7-80-9A-60	not found	not found
west entrance	00-22-90-5E-66-F0	A-09-E-102	1st floor West MSc
main entrance	00-22-90-38-7F-D0	not found	not found

Table 12.2: The AP a device connects with when entering the Faculty of Architecture

The AP a device connects with when entering the building via the West entrance, E-102', can also be found in Figure 12.5. Though it does not stand out compared to other APs. The MAC addresses of the APs the device connects with when entering the building via the East or Main entrance are not found, meaning the APs are not located on the map or listed in the table of APs. This implies it is not possible to relate the results to the data.

12.1.6. Recommendation

The results and conducted experiments has shown it is not possible to clearly find what, when and how frequent entrances of the Faculty of Architecture are used. The first and most important reason for that, is the time interval of approximately 5 minutes in which the system is recording. A person could be anywhere in the building at the moment of recording. A smaller time interval between the moments of recording would help in finding answers to the questions regarding the use of the entrances. Also, the existing layout of APs in the Faculty of Architecture is currently not designed for any other purpose than allowing a Wi-Fi compliant device to connect with the wireless eduroam network. Locating APs near the entrances of a building might help. Moreover, the fact the Faculty of Architecture has multiple entrances, in combination with the large time interval of recording, is what makes identification of the entrances difficult.

12.2. Association rules

The following section describes how movement patterns can be derived on building level, without considering the direction or order of the movement. An association rule mining algorithm (Agrawal, Imielinski, and Swami 1993) was used to identify groups of buildings that frequently visited in combination with each other. Firstly the algorithm is described briefly, then the results are presented and recommendation is given.

12.2.1. Association rules mining

Association rule mining is a technique to analyse what variables or items are commonly associated with each other in large databases. Probably the one of the main application is to analyse which items are commonly bought together by customers of a supermarket. As an example for this use case is an association rule of an itemset {bread, butter}, tells that in 80% of those transactions including {bread, butter}, also {milk} was present. In other words, 80% of the people who buy bread and butter also buy milk (Agrawal, Imielinski, and Swami 1993). Compared to sequence mining, association rule mining does not consider the order of items neither within, nor across transactions.

Thus every rule is composed by two itemsets, the *antecedent* {bread,butter} on the left-hand side, and the *consequent* {milk} on the right-hand side. The rule is denoted as {bread, butter} => {milk}.

12.2.2. Assosication rules of buildings

When a trajectory is simplified into a set of distinct buildings that the person visited, association rules for buildings can be derived. In this case the rule describes the set of buildings, or buildingset, that are commonly visited in combination. For example the rule {BK_City, Aula} => {Library} tells that a group of people who visited the buildings BK_City and Aula also visited the Library.

As association rule mining does not consider the order of buildings, nor the time spent in a building, it is important that these variables are appropriately handled and noise is filtered out prior running the algorithm.

In the first version the buildingsets were stored in a table as below, where the field *mac* contains the mac-address of a device and each remaining field represents a building. Value 1 is given if the device was recorded in a building, otherwise no value is given. This binary encoding is rather simplistic as it does not consider the amount of time spent in a building and therefore it does not allow to differentiate between occasional or regular visits.

mac	aula	bk_city	bouwcampus	btud	ctig	...
A	1	1			1	
B			1		1	
C		1				1
D	1					
E	1		1			

Table 12.3: uncategorized buildingset table

Therefore in the second version a distinction between *occasional*, *regular* and *frequent* stays was added to the buildingsets. The division between the categories is based on the 40 hour workweek and 1.5 hour lecture durations (see ??).

Category	hours/week	ID
occasional	≤ 0.5	1
regular	$> 0.5, \leq 5$	2
frequent	> 5	3

Table 12.4: Stay duration categories

The trajectories of approximately 14,000 devices were used to create the first set of association rules with categorized stay duration. At this stage only the noise was filtered from the data but not the stationary devices, and people carrying two devices were not accounted for. The time range of trajectories spanned from 31.03.2016 to 02.05.2016, approximately one month.

Although there are several measures to evaluate the interestingness of an association rule (Zhang et al. 2009), only *support* and *confidence* were used for testing purposes.

Support “The support for a rule is defined to be the fraction of transaction in the dataset that satisfy the union of items in the consequent and antecedent of the rule.” (Agrawal, Imieliński, and Swami 1993). In case of the rule {BK_City, Aula} => {Library}, the support is the percentage of the total dataset that includes BK_City, Aula and Library.

Confidence Confidence measures the strength of the rule, and is considered as a conditional probability. In case of the rule {BK_City, Aula} => {Library}, the confidence is the probability that Library is in the trajectory if both BK_City and Aula are in the trajectory (Agrawal, Imieliński, and Swami 1993; Anbukkarasy and Sairam 2013).

The most interesting rules are displayed in ??:

Supp	Conf	Covr	Strg	Lift	Levr	Antecedent	Consequent
0.02	0.86	0.02	6.92	6.51	0.01	dreibbelweg=2, ewi_lb=2 →	ewi_hb=2
0.01	0.74	0.01	24.80	3.38	0.00	btud=2, dreibbelweg=2, tbtm=2 →	cttg=2
0.01	0.70	0.01	30.21	3.21	0.00	aula=2, lr=2, ocp_io=2 →	cttg=2
0.01	0.70	0.01	29.22	2.44	0.00	aula=2, lr=2, ogz=2 →	btud=2
0.01	0.72	0.02	6.60	5.49	0.01	btud=2, ewi_lb=2 →	ewi_hb=2
0.01	0.73	0.01	15.17	5.52	0.01	aula=2, btud=2, ewi_lb=2 →	ewi_hb=2
0.01	0.72	0.01	16.24	5.44	0.00	btud=2, ctgg=2, ewi_lb=2 →	ewi_hb=2
0.01	0.90	0.01	17.31	6.81	0.01	btud=2, dreibbelweg=2, ewi_lb=2 →	ewi_hb=2
0.01	0.85	0.01	20.16	6.43	0.00	cttg=2, dreibbelweg=2, ewi_lb=2 →	ewi_hb=2
0.01	0.74	0.01	10.08	5.59	0.01	ewi_lb=2, tnw_tn=2 →	ewi_hb=2
0.01	0.86	0.01	21.85	6.51	0.00	dreibbelweg=2, ewi_lb=2, tnw_tn=2 →	ewi_hb=2
0.01	0.72	0.01	25.66	2.52	0.00	aula=1, tbtm=2 →	btud=2

Figure 12.7: Building set

In the buildingset of approx. 14,000 devices 2% was recorded in all of the buildings *Dreibbelweg, EWI-LB, EWI-HB* (Support = 0.02). There is an 86% chance that if a device is recorded in the buildings *Dreibbelweg, EWI-LB*, then it is also recorded in *EWI-HB* (Confidence = 0.86). And they spent on average between half hour to five hours a week in each building (dreibbelweg=2, ewi_lb=2, ewi_hb=2).

12.2.3. Recommendation

Association rule mining is a suitable technique to analize the occupancy of a group of buildings, but it is less suitable for analizing movement patterns. Therefore it is not handled more intensively in this project. However, this technique can potentially answer questions such as,

Which are the most visited buildings?,

Which buildings are islands?,

If a group of people vist building A, how likely that they will also visit building B?, All the people who visit building A, what other buildings do they visit as well?.

12.3. Distinguishing user groups

12.4. Occupancy

12.5. AP system

The setup of the system that logs the devices connected to access points is directly connected to the accuracy of the processed data. Currently, the APs register every device that is connected to it and the logging system receives all connected devices approximately every five minutes. Additionally, all access points are located indoors, logging every device carried by people using that building. These two aspects of the AP system limit the accuracy of the processed data and thus the movement patterns that can be derived.

Because the system logs every connected devices once every five minutes, a device will only be registered if the devices is connected for at least five minutes (not really true). This will result in discrepancies in the processed data. Devices and thus people walking by an AP will probably not be registered, for they are not connected to that AP for at least five minutes. This is unfortunate, because a person can travel a rather long distance in five minutes, e.g. making it hard to track people indoors. If the system would be logging every device all the time, irrespective of the time the device is connected, the tracking data would contain every AP that a device would connect to and thus provide much more accurate tracking data. Understandably, logging every user every second would result in huge amounts of data, which would most definitely result in perfor-

mance issues.

Secondly, because all scanners are located inside buildings, there is little to no information on people when they move from one building to another. Surely something can be told from the time it takes a device from the last scan in one building, to the first scan in the second building. But for outdoor tracking purposes, this system is limited. From some experiments that were conducted on the TU Delft Campus it can be concluded that a device located outdoors near a building can be detected by APs inside the building, but this depends on the antenna in the devices and the exact location of the device in respect to the AP. If more detailed information about movement outdoors is desired, it would be wise to also include outdoor APs in the system.

To improve further research, it is recommended to take the system of APs into account before actually conducting the research. If outdoor movement tracking is desired, outdoor APs are required. And if tracking indoors is one of the goals, the frequency of logging should be set to an interval that is in the order of magnitude of 10 seconds to one minute, taking data size in consideration.

12.6. Data reasoning

During this project a lot of data is handled. With all the data available and the processing to derive movement patterns one could ask: 'How reliable is the data?' and 'How accurately can we derive these movement patterns?'. Determining the working of the system of APs as described in ??, was a great step towards a reliable outcome. Knowing how the system works helped improve the processing steps that were taken, because the systems flaws could be taken into account and avoided. Additionally, when the first movement patterns were derived, common knowledge and knowledge about the TU Delft campus and its layout helped in validating these patterns.

Because the working of the system of APs is known, the dataset can be improved by filtering out people that are only registered for less than five minutes at one AP, indicating that they only were only passing by. This means that the states derived from the data are actually stay places of an individual. Another perspective could be that exactly those people that are only passing by are valuable for the dataset. When an individual is registered at four consecutive APs and each scan was less than five minutes, it can be concluded that this person is moving between those four APs. However, the current set-up of the AP system is not suitable enough to use only devices that have a session duration of less than five minutes. This would only work when the frequency of loggin is increased.

Moreover, the knowledge acquired from previous courses in the Geomatics programme and common knowledge about buildings and the TU Delft campus can be used to validate certain outcomes of the data processing. For example, it would seems very illogical that an individual could travel from Architecture to Aerospace Engineering and then to Industrial Design in five minutes. Such requirements could improve the final outcomes. This kind of reasoning became even more useful when zooming in to spatial level buildingpart. Using the knowledge of the building layout of Architecture, it could be concluded that moving from one floor to another is impossible without using one of the staircases. Such a conclusion could then be included in the processing, e.g. validating only movement between floors if one of the staircases is used.

For future research, it is desirable to use a higher frequency for logging the connected devices. This will ensure that a device is always registered and that its movement can be easily identified. Furthermore,

12.7. Visual exploration

13

Acknowledgements

We would like to take the opportunity to express our gratitude and regards to everyone that contributed to this project.

First, we would like to thank Edward Verbree, our supervisor, for the feedback provided during the course of the project. Additionally, we would like to thank Wilko Quack for his patience and support with all the issues we encountered when using the database.

We also would like to thank Bart Valks and Iljoesja Berdowski from the department of Facility Management and Real Estate for their feedback and guidance throughout the project and the challenges they proposed.

Additionally we would like to thank Jorge Gill for his feedback on our visualizations and his help for creating even better visualizations.

14

Appendix B




Figure 14.1: Movement on during weekdays and weekends at the BK




Figure 14.2: Movement on during weekdays at the BK




Figure 14.3: Movement on during weekends at the BK




Figure 14.4: Movement of mobile and static devices at the BK




Figure 14.5: Movement of mobile devices at the BK




Figure 14.6: Movement of static devices at the BK




Figure 14.7: Movement to and from the Bouwpub at the BK




Figure 14.8: Movement to the Bouwpub at the BK




Figure 14.9: Movement to and from the Restaurant at the BK




Figure 14.10: Movement to the Restaurant at the BK

Bibliography

- Agrawal, Rakesh, Tomasz Imieliński, and Arun Swami (1993). "Mining association rules between sets of items in large databases". In: *ACM SIGMOD Record*. Vol. 22. ACM, pp. 207–216. (Visited on 04/09/2015).
- Agrawal, Rakesh and Ramakrishnan Srikant (1994). "Fast algorithms for mining association rules". In: *Proc. 20th int. conf. very large data bases, VLDB*. Vol. 1215, pp. 487–499. URL: https://www.it.uu.se/edu/course/homepage/infoutv/ht08/vldb94_rj.pdf (visited on 04/09/2015).
- Anbukkarasy, G. and N. Sairam (2013). "Interesting Metrics Based Adaptive Prediction Technique for Knowledge Discovery". In: *International Journal of Engineering and Technology* 5.3, pp. 2069–2076. (Visited on 05/16/2016).
- Dodge, Somayeh, Robert Weibel, and Anna-Katharina Lautenschütz (2008). "Towards a taxonomy of movement patterns". In: *Information visualization* 7.3-4, pp. 240–252.
- Hunter, J. D. (2007). "Matplotlib: A 2D graphics environment". In: *Computing In Science & Engineering* 9.3, pp. 90–95.
- Mautz, Rainer (2012). "Indoor positioning technologies". PhD thesis. Habilitationsschrift ETH Zürich, 2012.
- Meneses, Filipe and Alberto Moreira (2012). "Large scale movement analysis from WiFi based location data". In: *Indoor Positioning and Indoor Navigation (IPIN), 2012 International Conference on*. IEEE, pp. 1–9.
- Parliament and Council of European Union (1995). *Directive (EC) 95/46/EC*. <http://eur-lex.europa.eu/LexUriServ/LexUriServ.do?uri=CELEX:31995L0046:en:HTML>.
- Pei, Jian et al. (2004). "Mining sequential patterns by pattern-growth: The prefixspan approach". In: *Knowledge and Data Engineering, IEEE Transactions on* 16.11, pp. 1424–1440.
- Radaelli, Laura et al. (2013). "Identifying Typical Movements among Indoor Objects—Concepts and Empirical Study". In: *Mobile Data Management (MDM), 2013 IEEE 14th International Conference on*. Vol. 1. IEEE, pp. 197–206.
- Van der Ham, M et al. (2014). "Rhythm of the campus". In:
- Zhang, Yuejin et al. (2009). "A Survey of Interestingness Measures for Association Rules". In: *International Conference on Business Intelligence and Financial Engineering, 2009. BIFE '09*. International Conference on Business Intelligence and Financial Engineering, 2009. BIFE '09. IEEE, pp. 460–463. DOI: 10.1109/BIFE.2009.110.
- Zhao, Shao et al. (2014). "Discovering People's Life Patterns from Anonymized WiFi Scanlists". In: *Ubiquitous Intelligence and Computing, 2014 IEEE 11th Intl Conf on and IEEE 11th Intl Conf on and Autonomic and Trusted Computing, and IEEE 14th Intl Conf on Scalable Computing and Communications and Its Associated Workshops (UTC-ATC-ScalCom)*. IEEE, pp. 276–283.